Agent Ops Pre-Flight Checklist

A short check to run before an AI agent, bot or scheduled pipeline is allowed to act on its own. It covers posting, deploying, emailing or spending money.

v1.0 — NeuralAlpha — neuralalphaai.substack.com

1. Identity: is it acting as the right account?

- ☐ The job checks *who it is* before the first write. For example, it calls the "who am I" endpoint and compares the result with a hard-coded expected handle or ID.
- ☐ A mismatch stops the run. It does not just log a warning.
- ☐ Credentials for different accounts live in separate files, and no job loads more than one set.
- ☐ You can name every account the job is able to write to.

2. Arming: can it fire when you think it is off?

- ☐ There is one obvious switch for "live" vs "dry run". The default is dry run.
- ☐ You have searched old schedulers, batch files and cron entries for jobs that are still armed. Old pipelines do not turn themselves off.
- ☐ Turning a job off is a change you can see, such as a commented line or a removed flag, with a dated backup.

3. Writes: what happens on the second attempt?

- ☐ Every outbound write records an "attempt" *before* sending, and a result (sent / failed / unknown) after.
- ☐ An "unknown" result (timeout, dropped connection) blocks the next send until a human looks. There is no automatic retry, because a retried post or payment is a duplicate.
- ☐ Each item has a stable ID, so you can tell whether it was already sent.
- ☐ A 200 response is treated as a claim. Where it matters, confirm by reading the object back.

4. Limits: what is the worst day?

- ☐ There is a hard cap per day (posts, emails, API calls), and the script enforces it rather than the prompt.
- ☐ The provider-side spend cap is set at the provider, not only in your code.
- ☐ Cost is logged per call, and there is an alert before the balance runs out, not after.
- ☐ Minimum spacing between actions is enforced.

5. Content: what can it say?

- ☐ Length, links and banned phrases are checked by code before sending.
- ☐ The rules on topics the agent must never touch (advice, claims, promises) are written down.
- ☐ Every "true story" the agent may tell is backed by a log, commit or receipt.

6. Deliverables: what does the customer actually get?

- ☐ You bought or downloaded your own product through the real checkout (in test mode) and opened the file.
- ☐ The delivered file's SHA-256 matches the build you meant to ship.
- ☐ The archive contains no internal notes, launch plans, keys or other customers' data.
- ☐ Builds are reproducible: rebuilding gives the same hash.

7. Environment: will it survive the night?

- ☐ The job has run once under the same scheduler, user, locale and encoding it uses in production. Windows cp932 vs UTF-8 is a classic overnight failure.
- ☐ Failures send a message to a person, not only to a log file.
- ☐ Someone other than the agent reads a summary of what it did, every day.

Rule of thumb: if a check above is "no", the agent runs in dry-run mode until it is "yes".

Every item comes from a real incident in a small automated publishing setup. Weekly notes on running AI agents in production: neuralalphaai.substack.com

Not financial, legal or security advice. Adapt to your own systems.